

# DBMS

## DBMS: Advanced Notes on Transactions and Related Concepts

This document provides advanced notes on Database Management Systems (DBMS), focusing extensively on transactions and their associated concepts. We'll delve beyond the basics, exploring nuanced aspects crucial for understanding and designing robust database applications.

### I. Transactions: Beyond ACID Properties

While the ACID properties (Atomicity, Consistency, Isolation, Durability) form the foundation of transactional integrity, a deeper understanding requires exploring their implications and limitations:

**Atomicity (All-or-Nothing):** A transaction is treated as a single, indivisible unit of work. Failure at any point necessitates complete rollback, leaving the database in its pre-transactional state. This involves sophisticated logging mechanisms (undo/redo logging) and recovery procedures to handle system crashes or failures. Advanced considerations include handling long-running transactions and partial rollbacks in specific scenarios (e.g., compensating transactions).

**Consistency (Maintaining Data Integrity):** A transaction must transform the database from one valid state to another. This relies heavily on constraints (primary keys, foreign keys, check constraints, etc.) and triggers that enforce business rules. Advanced topics include:

**Constraint Management:** Efficient constraint checking algorithms and optimization techniques.

**Multi-version Concurrency Control (MVCC):** Allows concurrent transactions to see consistent snapshots of the database, improving concurrency without sacrificing consistency.

**Serializability:** Ensuring that the concurrent execution of transactions produces the same result as if they were executed serially in some order.

Isolation (Preventing Interference): Multiple concurrent transactions should appear to execute in isolation, preventing interference and ensuring data integrity. Isolation levels (e.g., Read Uncommitted, Read Committed, Repeatable Read, Serializable) control the degree of isolation, trading off concurrency for consistency. Advanced concepts include:

Phantom Reads: A transaction reading data multiple times sees new rows inserted by another transaction.

Non-repeatable Reads: A transaction reading the same data multiple times sees changes made by another transaction.

Lost Updates: Two transactions update the same data, and one update is lost.

Deadlocks: Two or more transactions are blocked indefinitely, waiting for each other to release resources. Deadlock detection and resolution mechanisms (e.g., timeout, rollback) are critical.

Durability (Permanent Changes): Once a transaction commits, its changes are permanently stored and survive system failures. This requires robust logging and recovery mechanisms, including:

Write-Ahead Logging (WAL): Ensuring that log records are written to stable storage before data is modified.

Recovery Techniques: Strategies for restoring the database to a consistent state after a crash, including undo and redo operations based on the log. Advanced recovery techniques handle partial failures and complex scenarios.

## II. Concurrency Control Mechanisms

Various techniques manage concurrent access to the database, ensuring consistency and isolation:

Locking: Exclusive and shared locks prevent concurrent access to data items. Different locking protocols (e.g., strict two-phase locking, optimistic locking) offer varying degrees of concurrency and

overhead. Advanced locking strategies address granularity (fine-grained vs. coarse-grained) and deadlock prevention.

Timestamp Ordering: Transactions are assigned timestamps, and conflicts are resolved based on timestamp order. This avoids deadlocks but can lead to cascading rollbacks.

Multi-version Concurrency Control (MVCC): Each transaction operates on its own version of the database, minimizing conflicts and improving concurrency. Advanced MVCC implementations optimize read and write performance.

### III. Transaction Management in Distributed Databases

Managing transactions across multiple database systems introduces further complexity:

Two-Phase Commit (2PC): A protocol for coordinating commits across multiple sites. Advanced considerations include handling failures and optimizing performance.

Three-Phase Commit (3PC): An improvement over 2PC aiming to reduce blocking in failure scenarios.

Distributed Transaction Managers (DTMs): Software components that manage distributed transactions, handling communication and coordination between different database systems.

### IV. Advanced Transactional Features

**Nested Transactions:** Transactions can be nested within other transactions, providing finer-grained control and error handling.

**Sagga Transactions:** Transactions that can be partially committed, allowing for rollback of specific sub-transactions.

**Long-running Transactions:** Techniques for managing transactions that span extended periods, often involving compensating transactions to handle partial failures.

## V. Recovery and Backup Strategies

**Full Backups:** Complete copies of the database.

**Incremental Backups:** Copies of only the changes since the